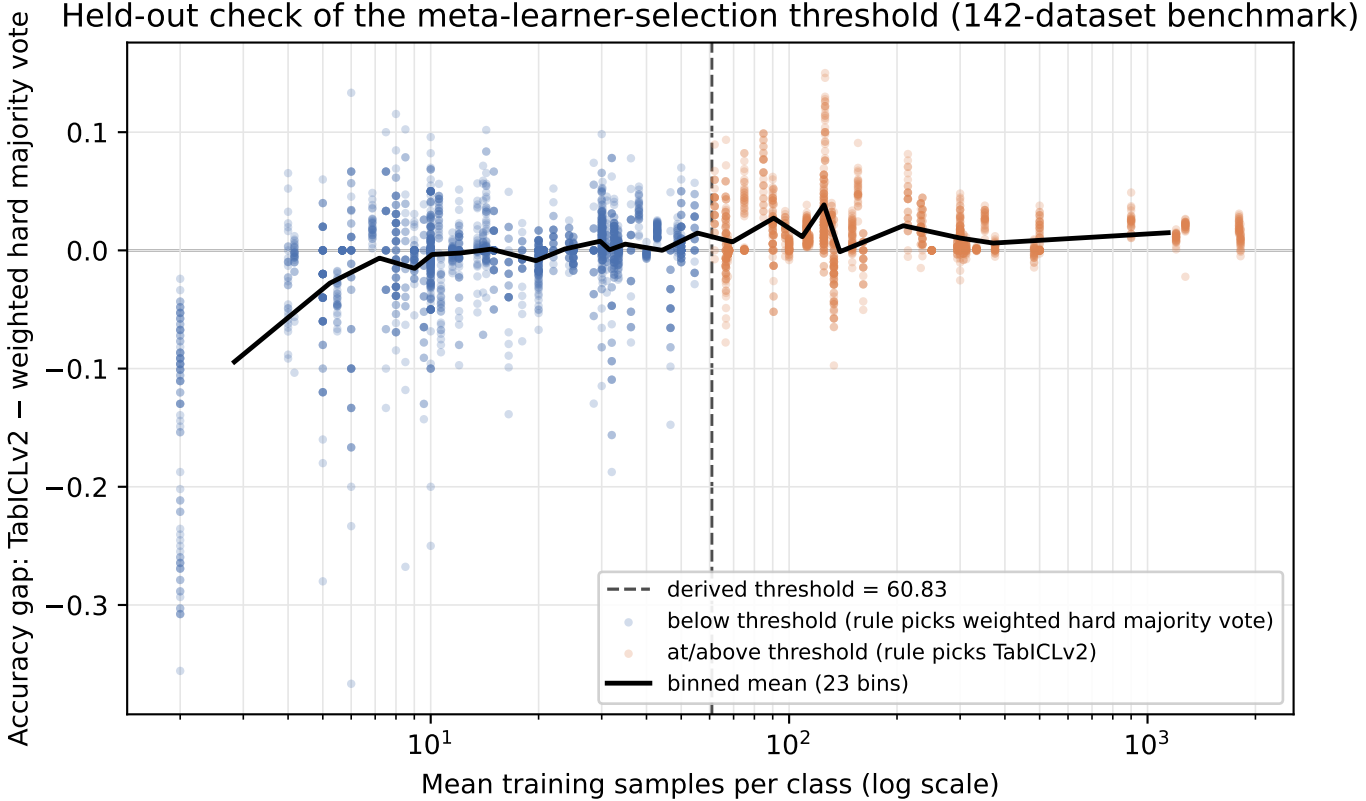
Held-out check of the meta-learner-selection threshold (142-dataset benchmark)
Accuracy gap: TabICLv2 − weighted hard majority vote
0.1
0.0
−0.1
−0.2
−0.3
10^1
10^2
10^3
Mean training samples per class (log scale)
derived threshold = 60.83
below threshold (rule picks weighted hard majority vote)
at/above threshold (rule picks TabICLv2)
binned mean (23 bins)